

---

# GRADIENT DESCENT AND MULTIPLE LINEAR REGRESSION

DR. FARHAD RAZAVI

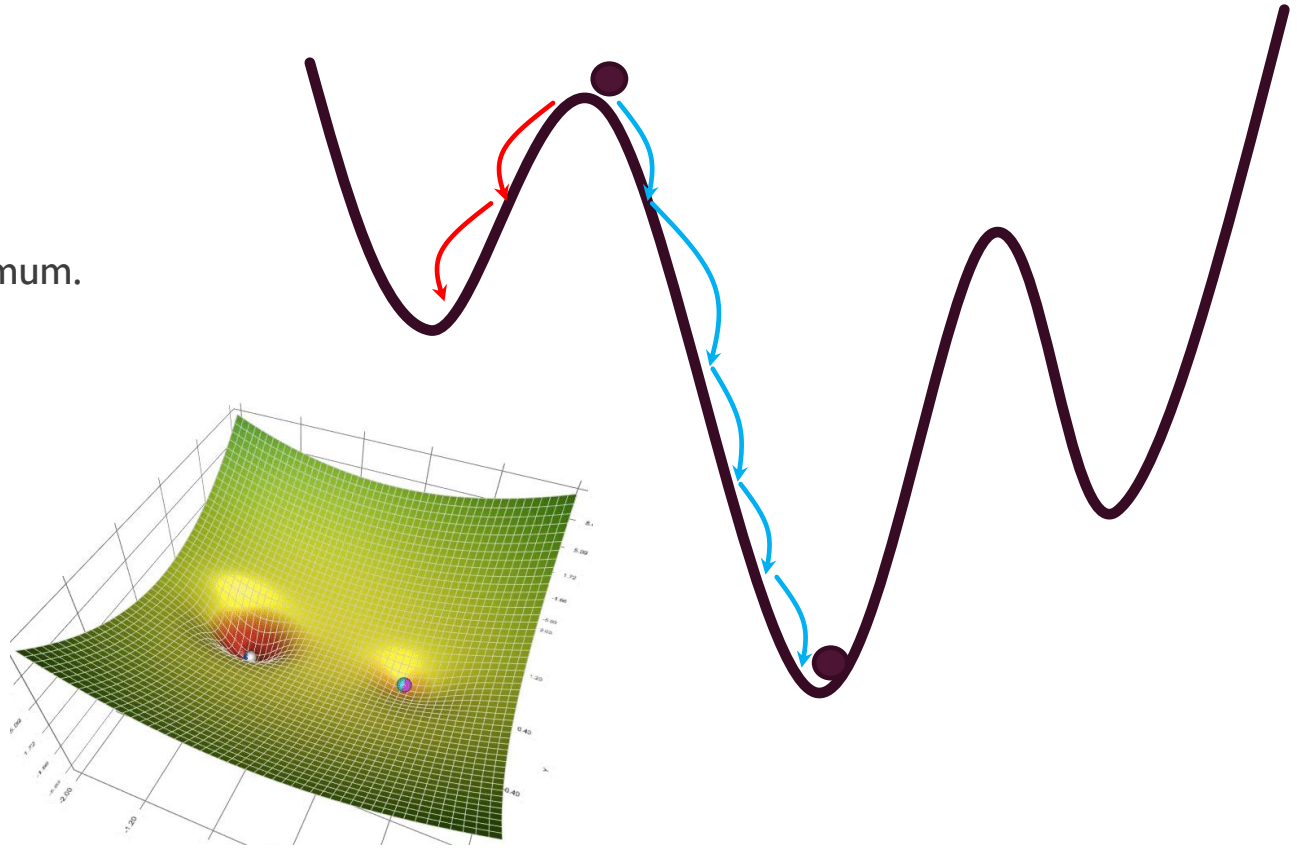


# OUTLINE

- Gradient Descent
- Linear regression for multiple features
- Gradient Descent for multiple features regression
- Feature engineering
- Polynomial regression

# GRADIENT DESCENT INTUITION

- Gradient Descent Algorithm:
  - Start with some  $w$  and  $b$ .
  - Keep changing  $w$ , and  $b$  to reduce  $J(w, b)$ .
  - Continue until we settle at or near a minimum.
- Drawbacks:
  - It could get stuck in a local minima.



# GRADIENT DESCENT FOR LINEAR REGRESSION

- Linear Regression:

- $f_{w,b}(x) = wx + b$

- $J(w, b) = \frac{1}{2m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)})^2$

- Gradient Descent

- Repeat until convergence {

$$w \rightarrow w - \alpha \frac{\partial}{\partial w} J(w, b) ;$$

$$b \rightarrow b - \alpha \frac{\partial}{\partial b} J(w, b);$$

}

$$\frac{\partial}{\partial w} J(w, b) = \frac{1}{m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)}) x^{(i)}$$

$$\frac{\partial}{\partial b} J(w, b) = \frac{1}{m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)})$$

- Cost function of Linear Regression is quadratic. This means it has only one minimum value, so we are guaranteed to reach global minimum.


# MULTIPLE FEATURES


Index $i$	Size in feet <sup>2</sup> $x_1$	Number of bedrooms $x_2$	Number of floors $x_3$	Age of home in years $x_4$	Price (\$) in \$1000's $y$
1	2104	5	1	45	460
2	1416	3	2	40	232
3	1534	3	2	30	315
...	...	...	...	...	...


- $x_j = j^{\text{th}}$  feature
- $n$  = number of features
- $\vec{x}^{(i)}$  = features of  $i^{\text{th}}$  training example
- $x_j^{(i)}$  value of feature  $j$  in  $i^{\text{th}}$  training example


# MODEL


- $f_{w,b}(x) = wx + b$
- $f_{w,b}(x) = w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + b$
- $f_{w,b}(x) = 0.1x_1 + 4x_2 + 10x_3 - 2x_4 + 80$ 

  
size

  
#bedrooms

  
#floors

  
years

  
base price

# VECTOR REPRESENTATION

- $f_{w,b}(x) = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$

$\vec{w} = [w_1 \ w_2 \ w_3 \ \dots \ w_n]$  parameters of the model

$b$  is a number

$\vec{x} = [x_1 \ x_2 \ x_3 \ \dots \ x_n]$  feature vector

- $f_{\vec{w},b}(\vec{x}) = \vec{w} \cdot \vec{x} + b = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$

↑  
dot product

↑  
Multiple Linear Regression (MLR)

# VECTORIZATION

- $f_{\vec{w},b}(\vec{x}) = \vec{w} \cdot \vec{x} + b = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$

```
w = np.array([0.1, 4.0, 10, -2])  
b = 4  
x = np.array([1513, 3, 2, 30])
```

```
f_wb = w[0] * x[0] +  
        w[1] * x[1] +  
        w[2] * x[2] + b
```

```
n = 4  
f_wb = 0  
for j in range(0, n):  
    f_wb = f_wb + w[j] * x[j]  
f_wb = f_wb + b
```

```
f_wb = np.dot(w, x) + b
```





# GRADIENT DESCENT FOR MLR

- Parameters:  $w_1, w_2, w_3, \dots, w_n$   
 $b$
- Models:  $f_{\vec{w}, b}(\vec{x}) = \vec{w} \cdot \vec{x} + b = w_1 x_1 + w_2 x_2 + \dots + w_n x_n + b$
- Cost function:  $J(w_1, w_2, w_3, \dots, w_n, b) = J(\vec{w}, b)$
- Gradient descent

- Repeat {

$$w_j = w_j - \alpha \frac{\partial}{\partial w_j} J(w_1, w_2, w_3, \dots, w_n, b) = w_j - \alpha \frac{1}{m} \sum_{i=1}^m (f_{w,b}(\vec{x}^{(i)}) - y^{(i)}) x_j^{(i)} \quad (j=1, \dots, n)$$

$$b = b - \alpha \frac{\partial}{\partial b} J(w_1, w_2, w_3, \dots, w_n, b) = b - \alpha \frac{1}{m} \sum_{i=1}^m (f_{w,b}(\vec{x}^{(i)}) - y^{(i)})$$

}

# CLOSED FORM SOLUTION

- There is a closed form solution for the “linear regression” parameters! It is called normal equation.

$$w = (XX^T)^{-1}Xy^T$$

- Disadvantages:
  - This is only available for Linear regression and does not generalize to other learning algorithms.
  - Slow when number of features are large ( > 10,000 )

# FEATURE SCALING

$$\widehat{price} = w_1 x_1 + w_2 x_2$$



size



#bedrooms

$x_1$ : size (feet<sup>2</sup>)

range: 300 – 2,000

large

$x_2$ : #bedrooms

range: 0 – 5

small

House:  $x_1 = 2000, x_2 = 5, price = \$500K$

one training example

size of the parameters  $w_1, w_2$  ?



$$w_1 = 50, w_2 = 0.1, b = 50$$

$$\widehat{price} = 50 * 2000 + 0.1 * 5 + 50$$





$$\widehat{price} = 100,050.5K$$


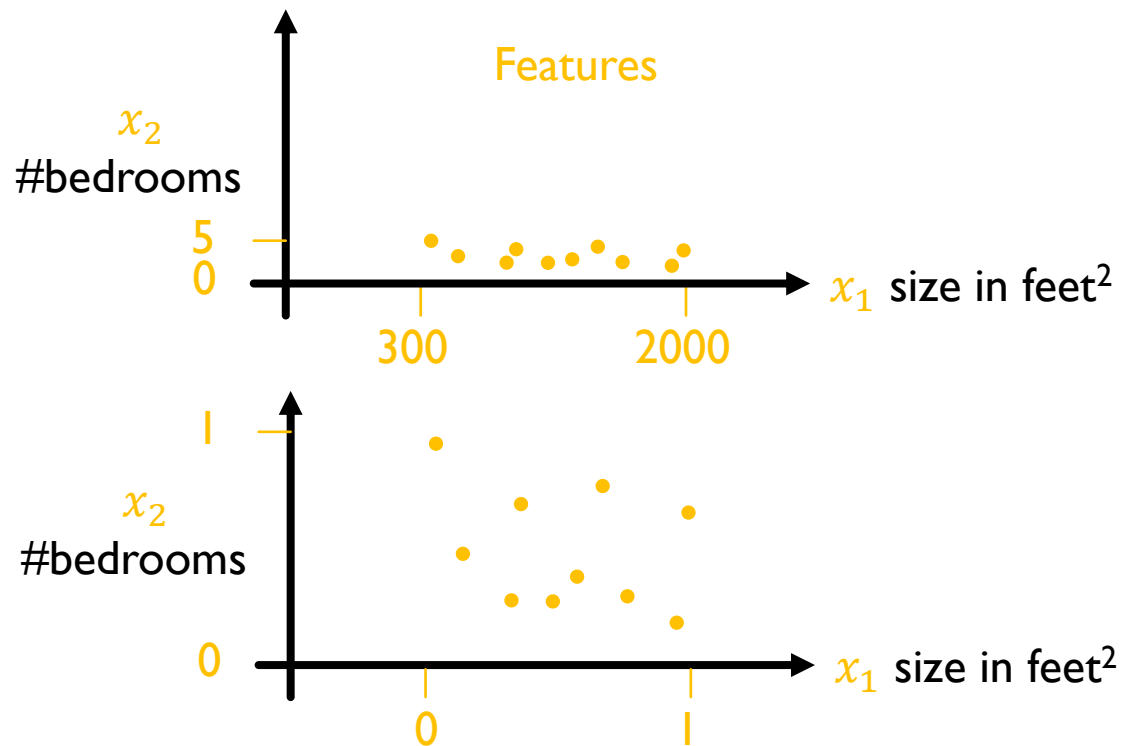
$$w_1 = 0.1, w_2 = 50, b = 50$$

$$\widehat{price} = 0.1 * 2000 + 50 * 5 + 50$$


$$\widehat{price} = 500K$$

# FEATURE SIZE AND PARAMETER SIZE

	Size of feature $x_j$	Size of parameter $w_j$
Size in feet <sup>2</sup>		
#bedrooms		




# FEATURE SCALING (MAX NORMALIZATION)



$$300 \leq x_1 \leq 2000$$

$$0 \leq x_2 \leq 5$$




$$x_{1,rescaled} = \frac{x_1}{2000}$$

$$0.15 \leq x_1 \leq 1$$

$$x_{2,rescaled} = \frac{x_2}{5}$$

$$0 \leq x_2 \leq 1$$

# FEATURE SCALING (MEAN NORMALIZATION)




$$300 \leq x_1 \leq 2000$$

$$\mu_1 = 600$$

$$0 \leq x_2 \leq 5$$

$$\mu_2 = 2.3$$




$$x_{1,rescaled} = \frac{x_1 - \mu_1}{2000 - 300}$$

$$-0.18 \leq x_1 \leq 0.82$$

$$x_{2,rescaled} = \frac{x_2 - \mu_2}{5 - 0}$$


$$-0.46 \leq x_1 \leq 0.54$$

# FEATURE SCALING (Z-SCORE NORMALIZATION)





$$300 \leq x_1 \leq 2000$$

$$\mu_1, \sigma_1 = 600, 450$$



$$0 \leq x_2 \leq 5$$

$$\mu_2, \sigma_2 = 2.3, 1.4$$



$$x_{1,rescaled} = \frac{x_1 - \mu_1}{\sigma_1}$$

$$-0.67 \leq x_1 \leq 3.1$$

$$x_{2,rescaled} = \frac{x_2 - \mu_2}{\sigma_2}$$

$$-1.6 \leq x_2 \leq 1.9$$


# FEATURE SCALING RULE OF THUMB

- Aim for about for  $-1 \leq x_j \leq 1$  for each feature  $x_j$
- Acceptable ranges  $\begin{cases} -3 \leq x_j \leq 3 \\ -0.3 \leq x_j \leq 0.3 \end{cases}$ 
  - $0 \leq x_j \leq 3$  Okay, no rescaling
  - $-2 \leq x_j \leq 0.5$  Okay, no rescaling
  - $-100 \leq x_j \leq 100$  Too large, rescale
  - $-0.001 \leq x_j \leq 0.001$  Too small, rescale
  - $98.6 \leq x_j \leq 105$  Too large, rescale



# CHECKING GRADIENT DESCENT

- Objective:  $\min_{\vec{w}, b} J(\vec{w}, b)$

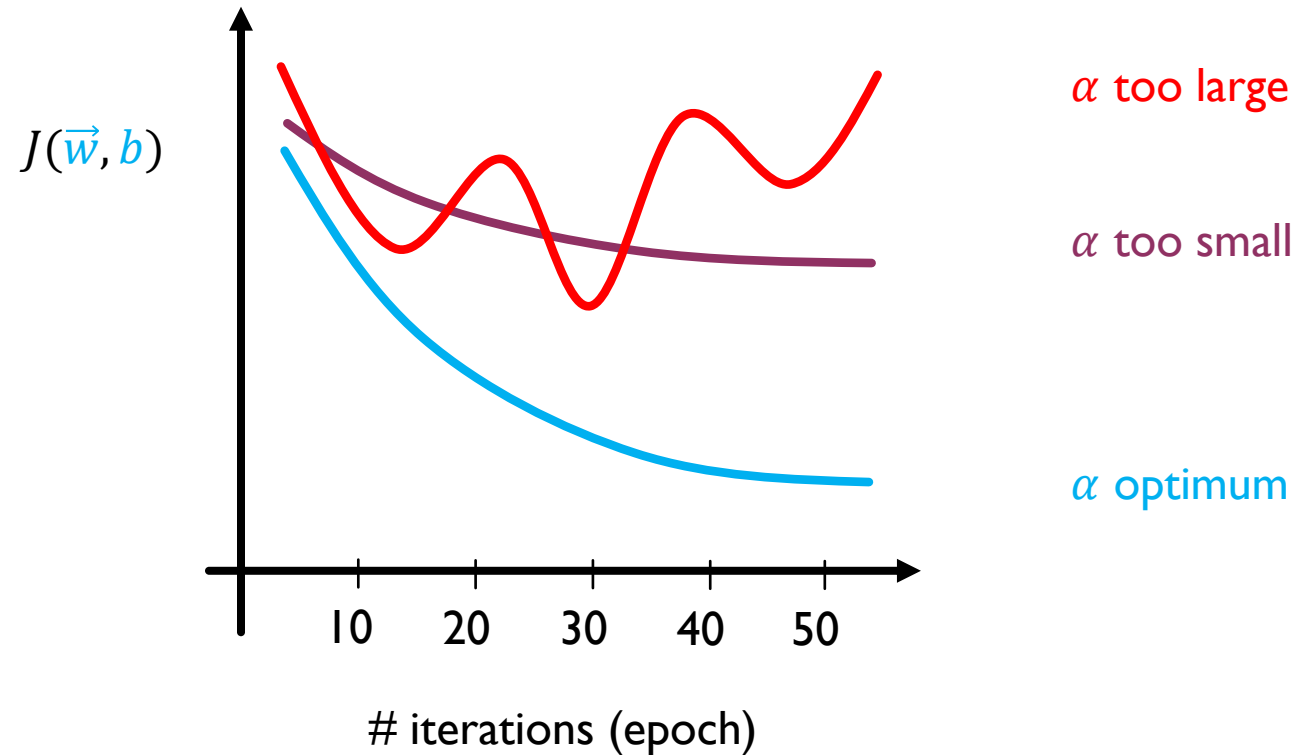


- If  $J(\vec{w}, b)$  is getting smaller on each iteration then GD is converging.
  - If you need to speed up the GD, try larger  $\alpha$  to see if the GD is still converging or not
- If  $J(\vec{w}, b)$  is getting bigger on each iteration then:
  - GD is diverging and your choice of  $\alpha$  was too big.
  - Maybe there is a bug in your code!
- The number of iteration is dependent on the application
  - Maybe 50, 1000, 100000 iteration is needed

# CHOOSING LEARNING RATE

- Values of  $\alpha$  to try:

0.001    0.003    0.01    0.03    0.1    0.3    1



# FEATURE ENGINEERING

$$f_{\vec{w},b}(\vec{x}) = w_1x_1 + w_2x_2 + b$$

↑                      ↙  
frontage              depth


$$area = frontage \times depth$$

$$x_3 = x_1 x_2 \quad \text{new feature}$$


$$f_{\vec{w},b}(\vec{x}) = w_1x_1 + w_2x_2 + w_3x_3 + b$$

- Feature engineering :

- Using intuition or knowledge to design new features, by transforming or combining original features.
- Needs human insight.



# POLYNOMIAL REGRESSION



$$f_{\vec{w},b}(x) = w_1x + b$$

$$f_{\vec{w},b}(x) = w_1x + w_2x^2 + b$$

$$f_{\vec{w},b}(x) = w_1x + w_2x^2 + w_3x^3 + b$$

↑            ↑            ↑  
size      size<sup>2</sup>    size<sup>3</sup>

Normalization becomes extremely important

# ACKNOWLEDGEMENT

- The material are based on Prof. Andrew Ng course on this subject.